



(A) Full Training Horizon: Convergence to Identical Ceiling



(B) The Masking Effect: 100k Dynamics vs. 2M Saturation



Takeaway: 2M timesteps masks critical real-world differences. At 100k steps, FP3O achieves superior autonomous safety (13.8% shield rate vs >32%), while avoiding representation interference. At 2M steps, all algorithms saturate near +16.71, concealing inductive bias and safety failures.